



Orchestrating Containers

Kubernetes Basics

Eueung Mulyana

<http://eueung.github.io/docker-stuff/kubernetes>

CodeLabs | Attribution-ShareAlike CC BY-SA

Outline

Kubernetes - Short Introduction

Parts & Components

Getting Started

`kubectl run + kubectl expose`



Kubernetes Short Introduction

What is Kubernetes?

The name Kubernetes originates from Greek, meaning "helmsman" or "pilot", and is the root of "governor" and "cybernetic".

K8s is an abbreviation derived by replacing the 8 letters "ubernete" with 8.

With Kubernetes you can deploy a full cluster of **multi-tiered** containers (frontend, backend, etc.) with a **single** configuration file and a **single** command (Ref).

Kubernetes is an open-source platform for **automating** deployment, scaling, and operations of **application containers** across **clusters** of hosts, providing **container-centric** infrastructure.

With Kubernetes, you are able to quickly and efficiently respond to customer demand:

- **Deploy** your applications quickly and predictably.
- **Scale** your applications on the fly.
- Seamlessly **roll out** new features.
- **Optimize** use of your hardware by using only the resources you need

Kubernetes is:

- **portable**: public, private, hybrid, multi-cloud
- **extensible**: modular, pluggable, hookable, composable
- **self-healing**: auto-placement, auto-restart, auto-replication, auto-scaling

Ref: kubernetes.io





Kubernetes

- Container orchestrator
- Runs and manages containers
- Supports multiple cloud and bare-metal environments
- Inspired and informed by Google's experiences and internal systems
- 100% Open source, written in Go
- Manage applications, not machines

Ref: [Kubernetes Intro and Update @thockin](#)

Architecture

A running Kubernetes cluster contains **node agents** (**kubelet**) and **master** components (**apiserver**, **scheduler**, etc), on top of a distributed storage solution.



K8s Node

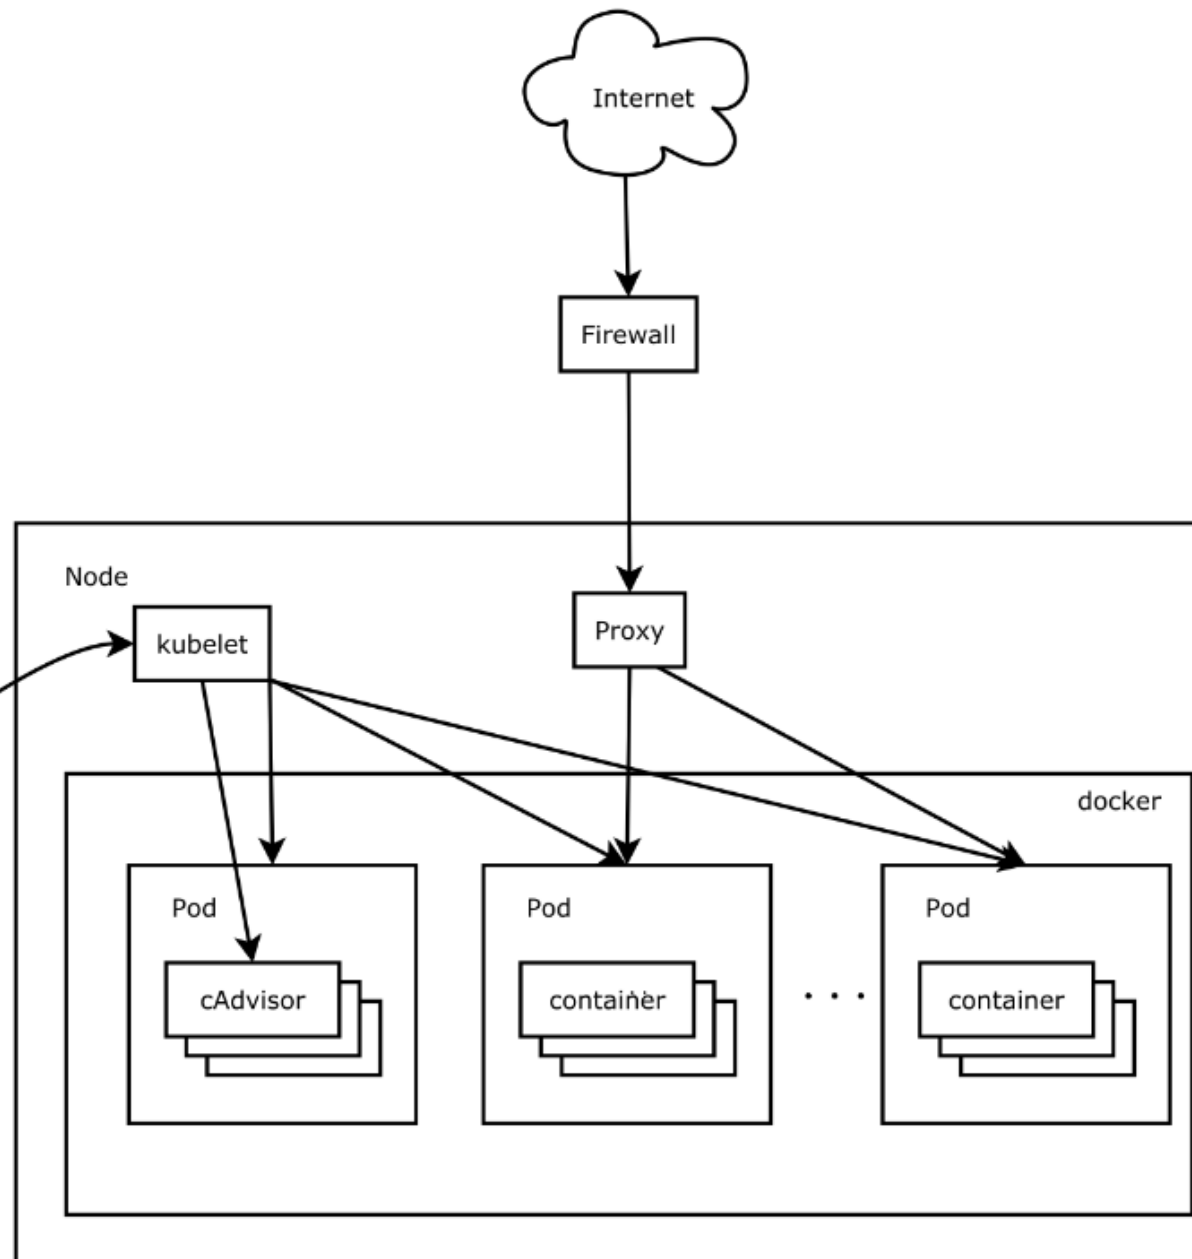
Worker / Minion

The Kubernetes node has the services necessary to **run** application containers and **be managed** from the master systems.

Each node runs container engine e.g. Docker. This engine takes care of the details of downloading images and running containers.

kubelet

kube-proxy



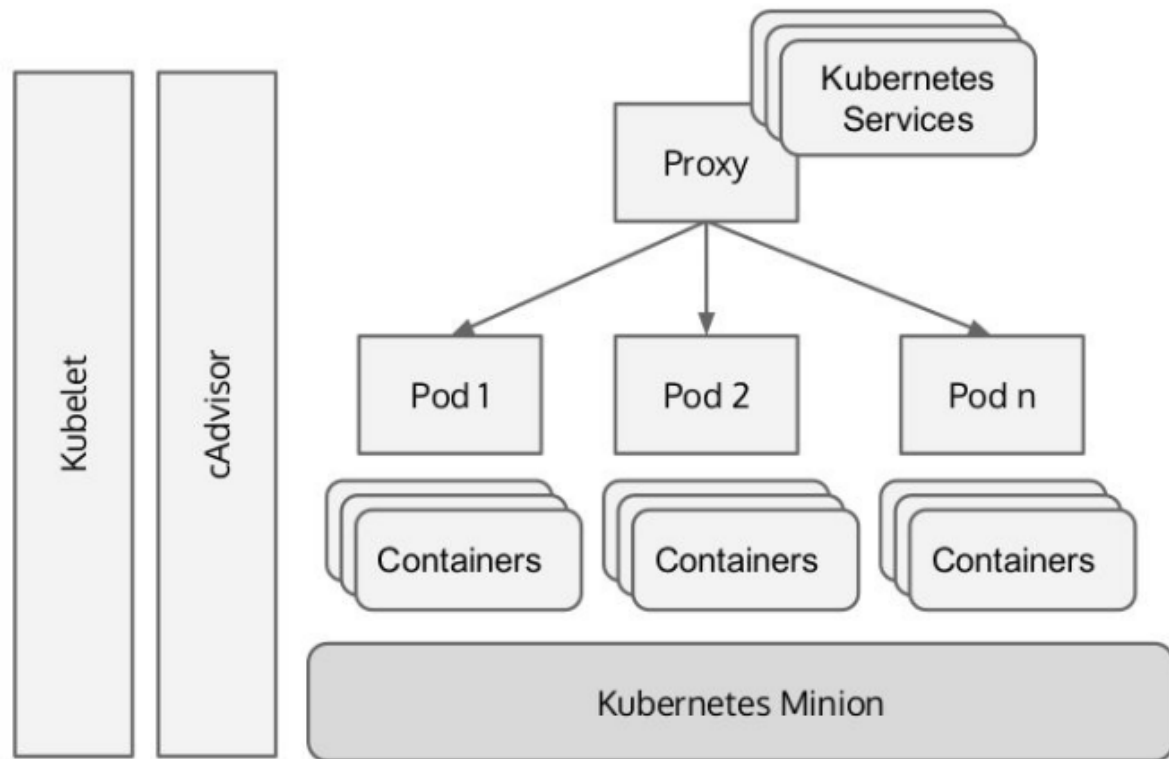
K8s Node

kubelet

The kubelet manages **pods** and their containers, their images, their volumes, etc.

kube-proxy

Each node also runs a simple network proxy and load balancer. This reflects **services** as defined in the Kubernetes API (apiserver) on each node and can do simple TCP and UDP stream forwarding (round robin) across a set of backends.



K8s Master

Control Plane

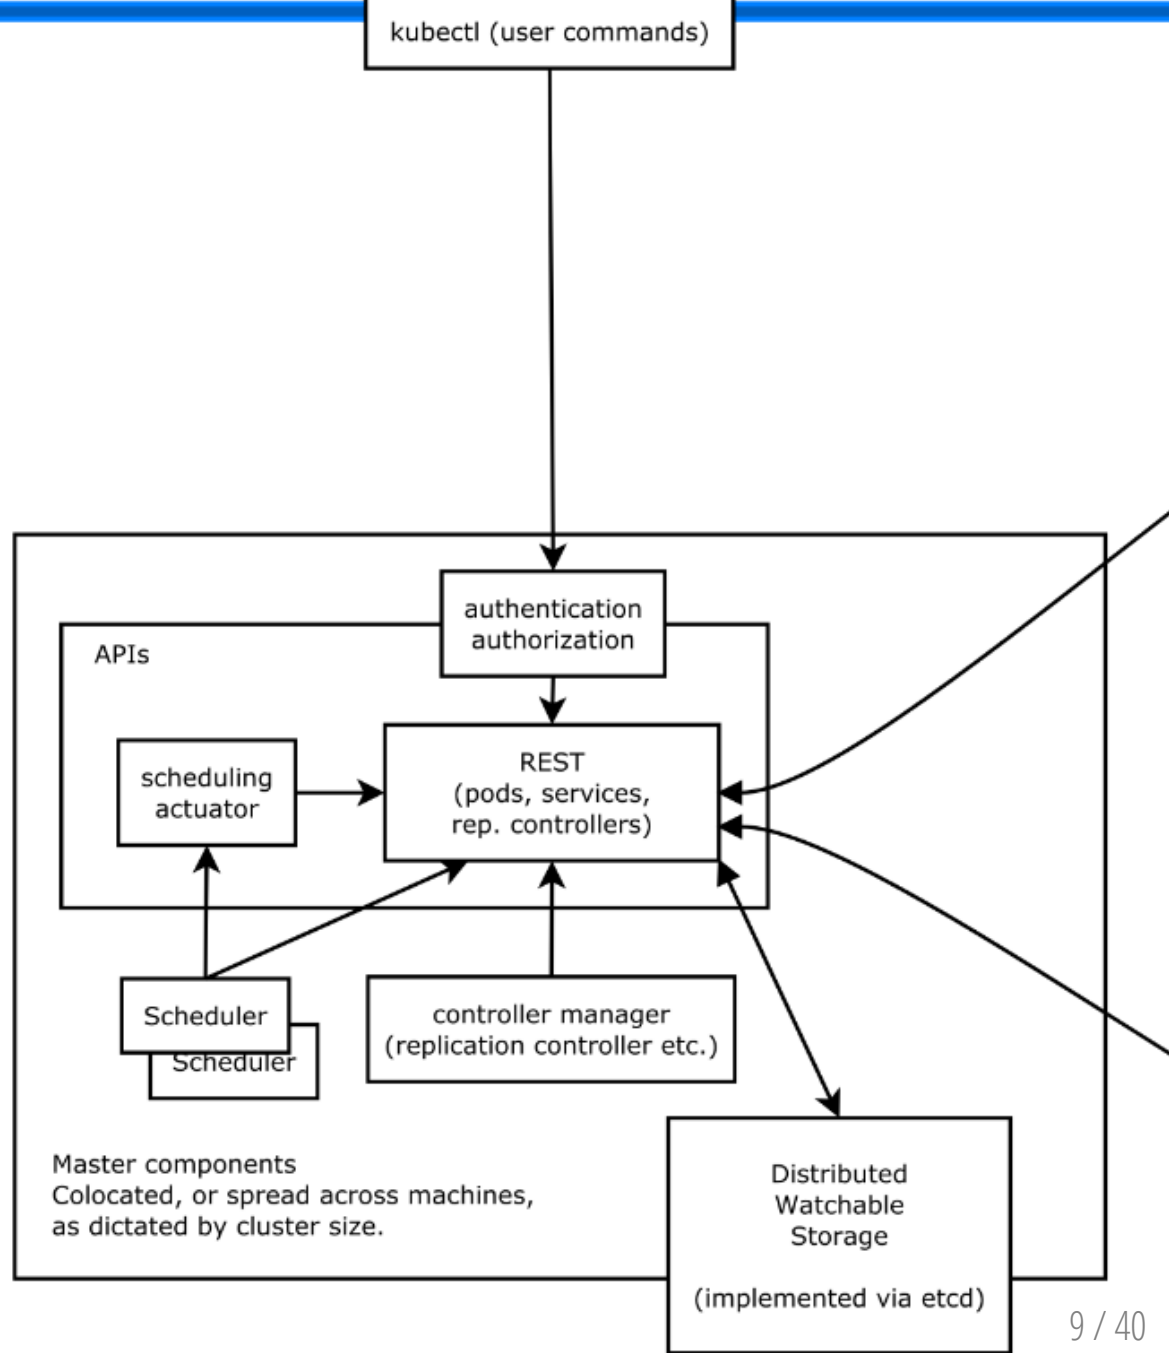
The Kubernetes control plane is split into a set of components. Currently they all run on a single master node. These components work together to provide a unified view of the cluster.

etcd

All persistent master state is stored in an instance of etcd.

This provides a great way to store configuration data reliably.

With watch support, coordinating components can be notified very quickly of changes.

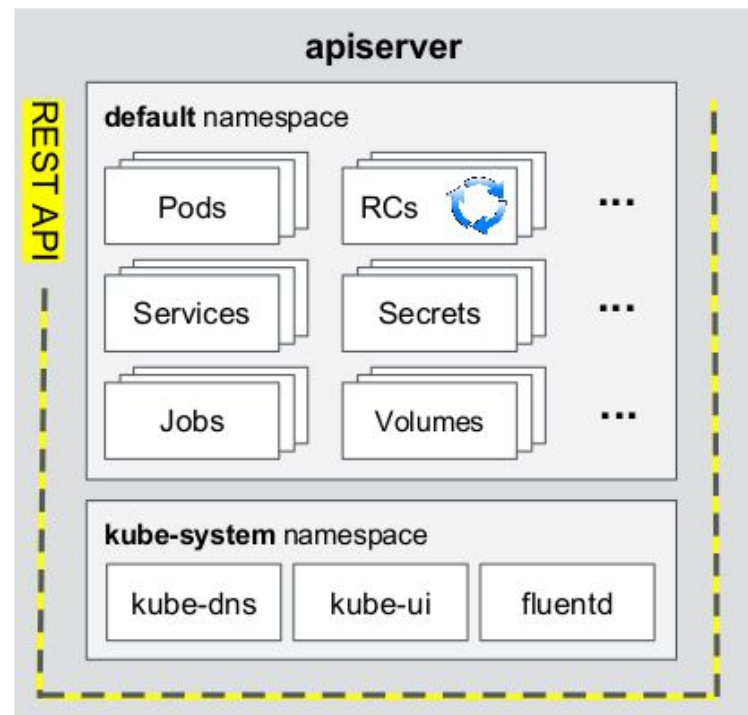
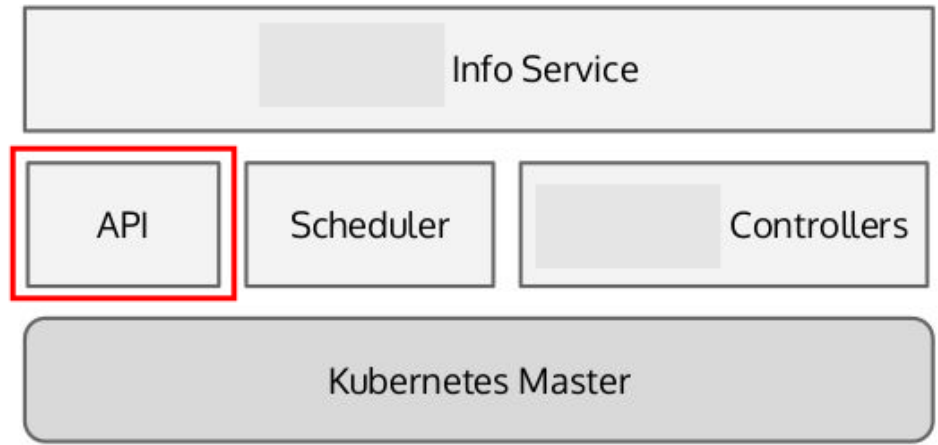


K8s Master

etcd

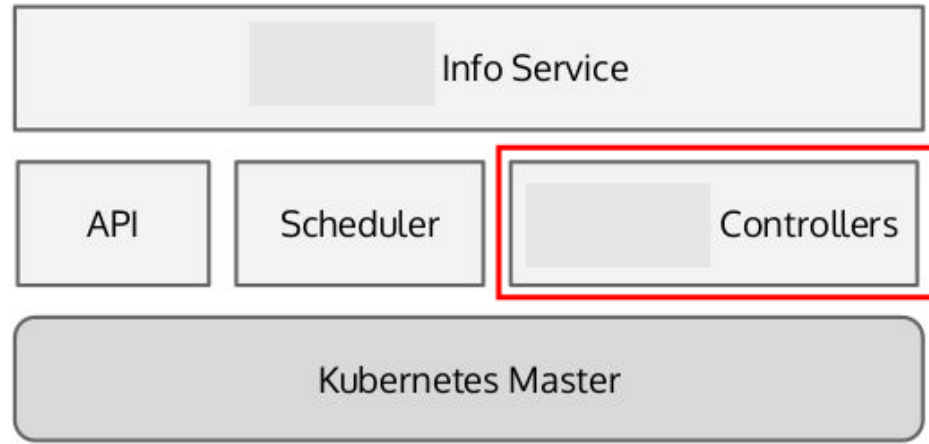
API Server

The apiserver serves up the Kubernetes API. It is intended to be a CRUD-y server, with most/all business logic implemented in separate components or in plug-ins. It mainly processes REST operations, validates them, and updates the corresponding objects in etcd (and eventually other stores).



Scheduler

The scheduler binds unscheduled pods to nodes via the **/binding** API. The scheduler is pluggable, support for multiple cluster schedulers and user-provided schedulers in the plan.

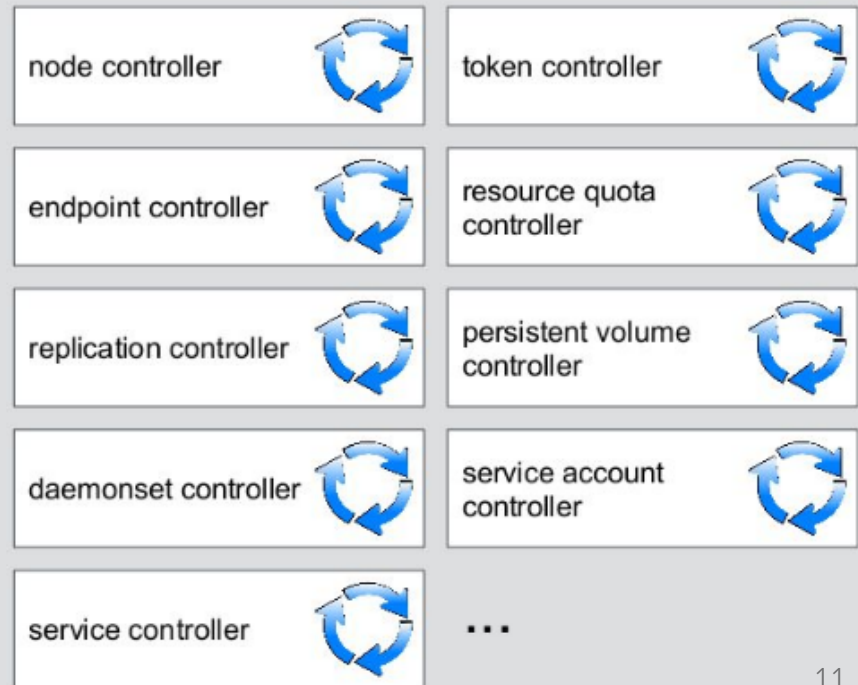


Controller Manager

All other cluster-level functions are currently performed by the Controller Manager.

For instance, Endpoints objects are created and updated by the **endpoints controller**, and nodes are discovered, managed, and monitored by the **node controller**. These could eventually be split into separate components to make them independently pluggable. The **replicationcontroller** is a mechanism that is layered on top of the simple pod API.

controller-manager





Kubernetes Parts & Components

Containers




Containers

Base Asset

Containers



Pods



Containers

Pods

A pod is a co-located group of containers ...

Volumes



Containers

Pods

Pods with **Volumes**

A pod is a co-located group of containers and volumes.

A volume is a directory, possibly with some data in it, which is accessible to a Container as part of its filesystem.

Kubernetes volumes build upon Docker Volumes, adding provisioning of the volume directory and/or device.

Labels



Containers

Pods

Pods with Volumes

Labels

A label is a key/value pair that is attached to a resource, such as a pod, to convey a user-defined identifying attribute.

Labels can be used to organize and to select subsets of resources.

Replication Controller



Containers

Pods

Pods with Volumes


Labels

Replication Controllers

A replication controller ensures that a specified number of pod replicas are running at any one time.

It both allows for easy scaling of replicated systems and handles re-creation of a pod when the machine it is on reboots or otherwise fails.

Replication Controller



Pods

Containers

Pods

Pods with Volumes

Labels

Replication Controllers

Creating labeled pods with a
(labeled) RC

Service



Service



Containers

Pods

Pods with Volumes

Labels


Replication Controllers

Services

Services & labeled Services

A service defines a set of pods and a means by which to access them, such as single stable IP address and corresponding DNS name.

Put it all together



Containers
Pods
Pods with Volumes
Labels
Replication Controllers
Services



Kubernetes Getting Started



node1 **RPI 3**

192.168.1.102

node2 **RPI 2**

192.168.1.101

node3 **RPI 2**

192.168.1.100

kubernetes-on-arm v0.7.0

by @luxas

```
$ dpkg -i kube-systemd.deb
$ kube-config install
#rpi or rpi-2, hypriotos

$ gzip -dc images.tar.gz | docker load

$ kube-config info
Architecture: armv7l
Kernel: Linux 4.1.17
CPU: 4 cores x 1200 MHz

Used RAM Memory: 111 MiB - RAM Memory: 925 MiB

Used disk space: 2.4GB (2424044 KB)
Free disk space: 11GB (11382472 KB)

SD Card/deb package was built: 21-03-2016 21:06

kubernetes-on-arm: Latest commit: 1d0bbe1 - Version: 0.7.0

systemd version: v215
docker version: v1.10.3
kubernetes client version: v1.2.0
```



\$ docker images

REPOSITORY	TAG	IMAGE ID	CREATED
kubernetesarm/etcd	0.7.0	3550c0c4f205	8 weeks ago
kubernetesarm/etcd	latest	3550c0c4f205	8 weeks ago
kubernetesarm/grafana	0.7.0	8ea26f4ef5a3	8 weeks ago
kubernetesarm/grafana	latest	8ea26f4ef5a3	8 weeks ago
kubernetesarm/influxdb	0.7.0	447951a687c6	8 weeks ago
kubernetesarm/influxdb	latest	447951a687c6	8 weeks ago
kubernetesarm/heapster	0.7.0	1d034674fc0e	8 weeks ago
kubernetesarm/heapster	latest	1d034674fc0e	8 weeks ago
kubernetesarm/loadbalancer	0.7.0	6b262f2fd318	8 weeks ago
kubernetesarm/loadbalancer	latest	6b262f2fd318	8 weeks ago
kubernetesarm/registry	0.7.0	4c9c964f89e9	8 weeks ago
kubernetesarm/registry	latest	4c9c964f89e9	8 weeks ago
kubernetesarm/exechealthz	0.7.0	57f77542ac5c	8 weeks ago
kubernetesarm/exechealthz	latest	57f77542ac5c	8 weeks ago
kubernetesarm/kube2sky	0.7.0	770cccac4236	8 weeks ago
kubernetesarm/kube2sky	latest	770cccac4236	8 weeks ago
kubernetesarm/skydns	0.7.0	30ea4958f939	8 weeks ago
kubernetesarm/skydns	latest	30ea4958f939	8 weeks ago
kubernetesarm/pause	0.7.0	05edc969256e	8 weeks ago
kubernetesarm/pause	latest	05edc969256e	8 weeks ago
kubernetesarm/hyperkube	0.7.0	380def049467	8 weeks ago
kubernetesarm/hyperkube	latest	380def049467	8 weeks ago
kubernetesarm/flannel	0.7.0	8a35c629399f	8 weeks ago
kubernetesarm/flannel	latest	8a35c629399f	8 weeks ago
hypriot/rpi-swarm	latest	c298de062190	10 weeks ago

Master

```
$ kube-config enable-master
```

```
Disabling k8s if it is running
```

```
Checks so all images are present
```

```
Transferring images to system-docker, if necessary
```

```
Copying kubernetesonarm/etcd to system-docker
```

```
...
```

```
Kubernetes master services enabled
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED
7bf736d84451	kubernetesonarm/hyperkube	"/hyperkube controlle"	2 minutes a
5560f06bcf47	kubernetesonarm/hyperkube	"/hyperkube proxy --m"	2 minutes a
0fb169567946	kubernetesonarm/hyperkube	"/hyperkube scheduler"	2 minutes a
11c03a40412c	kubernetesonarm/hyperkube	"/hyperkube apiserver"	2 minutes a
ede858b35dfa	kubernetesonarm/pause	"/pause"	2 minutes a
5e38dac4fd19	kubernetesonarm/hyperkube	"/hyperkube kubelet -"	3 minutes a

Worker

```
$ kube-config enable-worker 192.168.1.102
```

```
Disabling k8s if it is running
```

```
Using master ip: 192.168.1.102
```

```
Checks so all images are present
```

```
Transferring images to system-docker, if necessary
```

```
Copying kubernetesonarm/flannel to system-docker
```

```
Created symlink from /etc/systemd/system/multi-user.target.wants/flannel.service to
```

```
Starting worker components in docker containers
```

```
Created symlink from /etc/systemd/system/multi-user.target.wants/k8s-worker.service
```

```
Kubernetes worker services enabled
```

Check

```
$ kubectl get nodes
```

NAME	STATUS	AGE
192.168.1.100	Ready	26s
192.168.1.101	Ready	7m
192.168.1.102	Ready	4h

```
$ kubectl cluster-info
```

```
Kubernetes master is running at http://localhost:8080
```

Test #1

```
$ kubectl run my-nginx --image=lucas/nginx-test --replicas=3 --expose --port=80
service "my-nginx" created
deployment "my-nginx" created
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
k8s-master-192.168.1.102	4/4	Running	1	5h
my-nginx-3795026575-9w8mw	0/1	ContainerCreating	0	14s
my-nginx-3795026575-miz3d	0/1	ContainerCreating	0	14s
my-nginx-3795026575-zy2d4	0/1	ContainerCreating	0	14s

```
$ kubectl get svc
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	10.0.0.1	<none>	443/TCP	5h
my-nginx	10.0.0.129	<none>	80/TCP	36s

```
$ kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
my-nginx	3	3	3	1	1m

```
$ curl 10.0.0.129
```

```
<p>WELCOME TO NGINX</p>
```

Test #2

```
$ kubectl run hello-kube --image=hypriot/rpi-nano-httpd --port=80
deployment "hello-kube" created
```

```
$ kubectl expose deployment hello-kube --type="LoadBalancer" --external-ip="192.168.
service "hello-kube" exposed
```

```
$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	NODE
hello-kube-1079346743-2knj3	0/1	ImagePullBackOff	0	55s	192.
k8s-master-192.168.1.102	4/4	Running	1	8h	192.

```
$ kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
hello-kube	1	1	1	1	1m

```
$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	NODE
hello-kube-1079346743-2knj3	1/1	Running	0	1m	192.168.1.100
k8s-master-192.168.1.102	4/4	Running	1	8h	192.168.1.102

```
$ kubectl get svc
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-kube	10.0.0.197	,192.168.1.102	80/TCP	44s
kubernetes	10.0.0.1	<none>	443/TCP	8h

```
$ curl 10.0.0.197
```

```
<html><head><title>Pi armed with Docker by Hypriot</title>
  <body style="width: 100%; background-color: black;">
    <div id="main" style="margin: 100px auto 0 auto; width: 800px;">
      
  </body></html>
```


Test #3

```
$ kubectl run hello-kube --image=hypriot/rpi-nano-httpd --replicas=2 --port=80
deployment "hello-kube" created
```

```
$ kubectl expose deployment hello-kube --type="LoadBalancer" --external-ip="192.168.1.102"
service "hello-kube" exposed
```

```
$ kubectl get svc
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-kube	10.0.0.221	,192.168.1.102	80/TCP	10s
kubernetes	10.0.0.1	<none>	443/TCP	8h

```
$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	NODE
hello-kube-1079346743-44yig	1/1	Running	0	1m	192.168.1.100
hello-kube-1079346743-upzxy	1/1	Running	0	1m	192.168.1.101
k8s-master-192.168.1.102	4/4	Running	1	8h	192.168.1.102

```
$ curl 10.0.0.221
```

```
$ curl 192.168.1.102
```

```
<html><head><title>Pi armed with Docker by Hypriot</title>
  <body style="width: 100%; background-color: black;">
    <div id="main" style="margin: 100px auto 0 auto; width: 800px;">
      
    </div>
  </body></html>
```

Test #4

```
$ kubectl run hello-kube --image=hypriot/rpi-nano-httpd --replicas=3 --port=80
deployment "hello-kube" created
```

```
$ kubectl expose deployment hello-kube --port=8300 --target-port=80 --type="LoadBalancing"
service "hello-kube" exposed
```

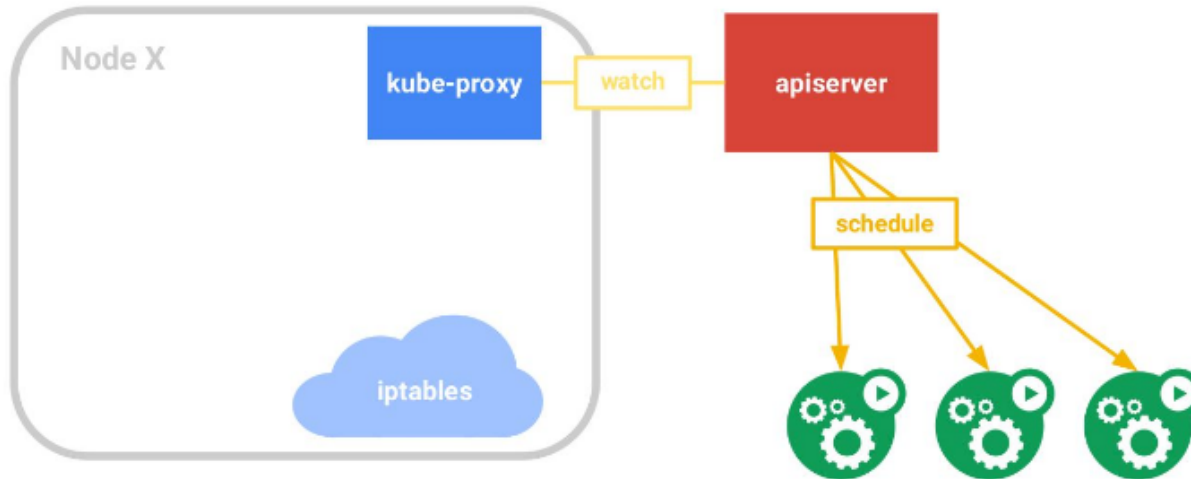

```
$ kubectl get svc
```

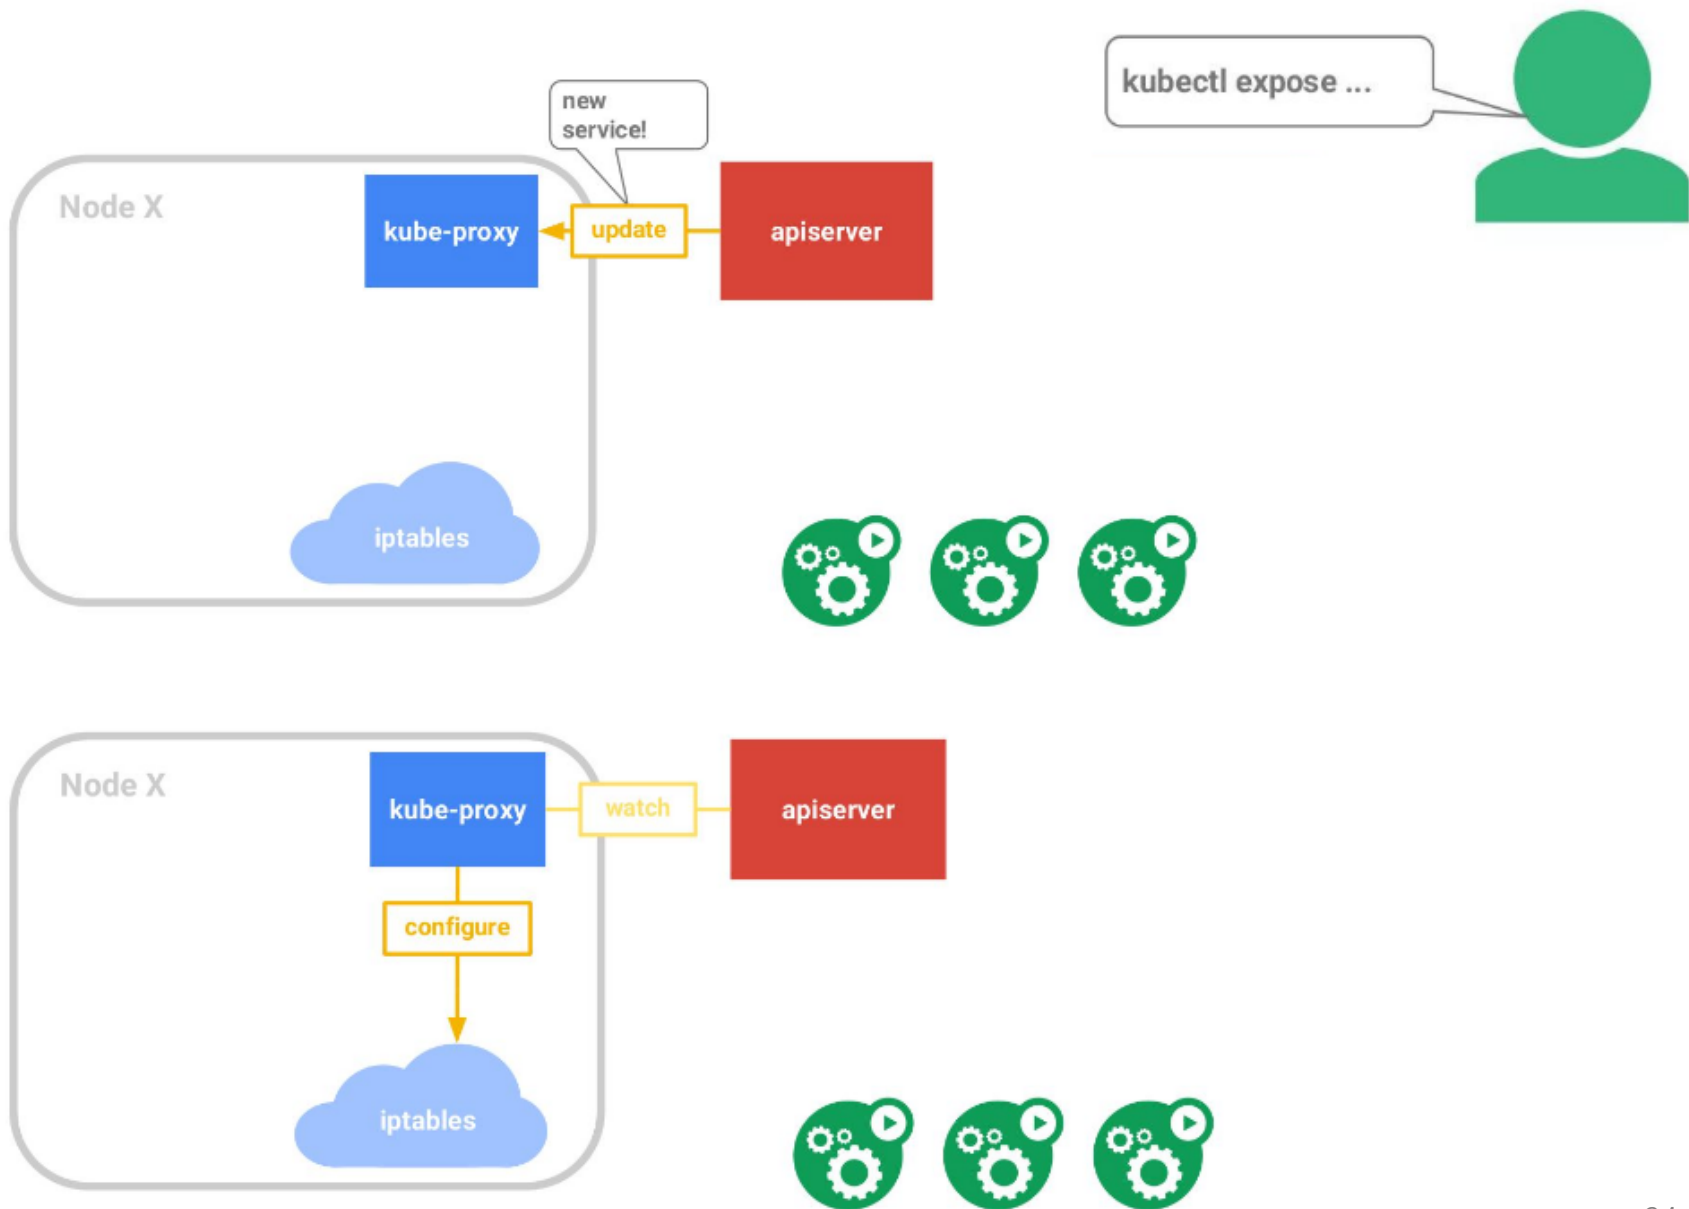
NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-kube	10.0.0.124	,192.168.1.102	8300/TCP	8s
kubernetes	10.0.0.1	<none>	443/TCP	7h


```
$ kubectl delete service,deployment hello-kube
```

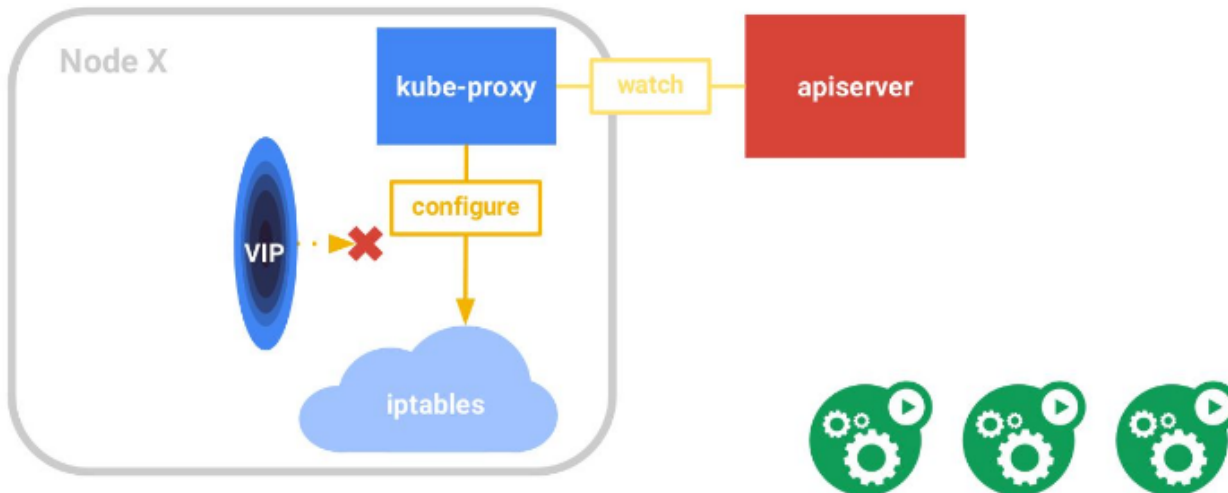
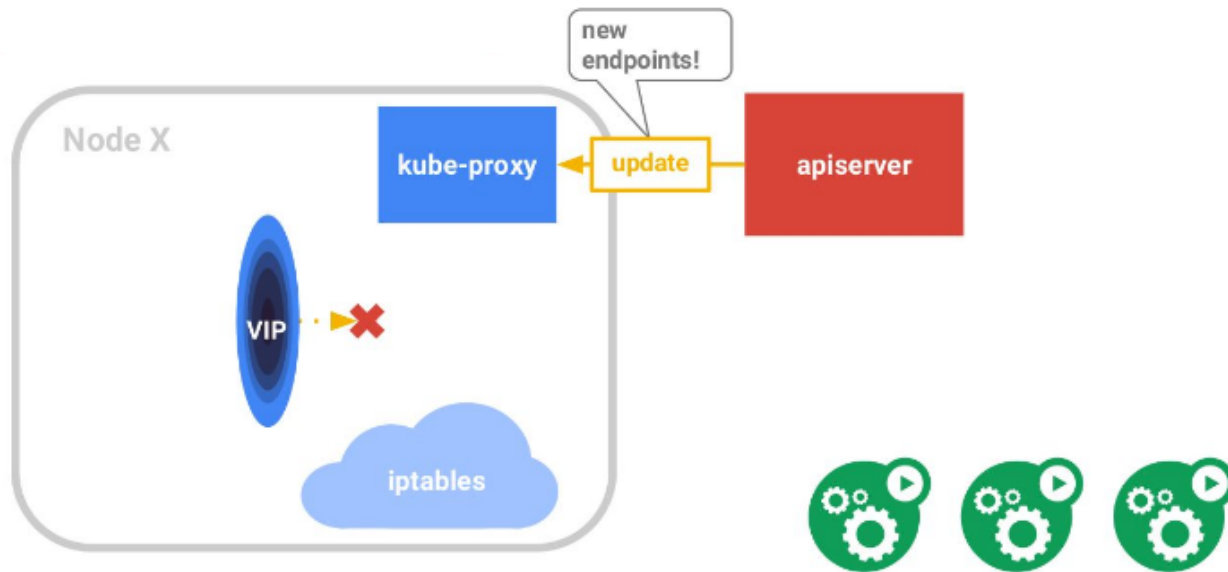


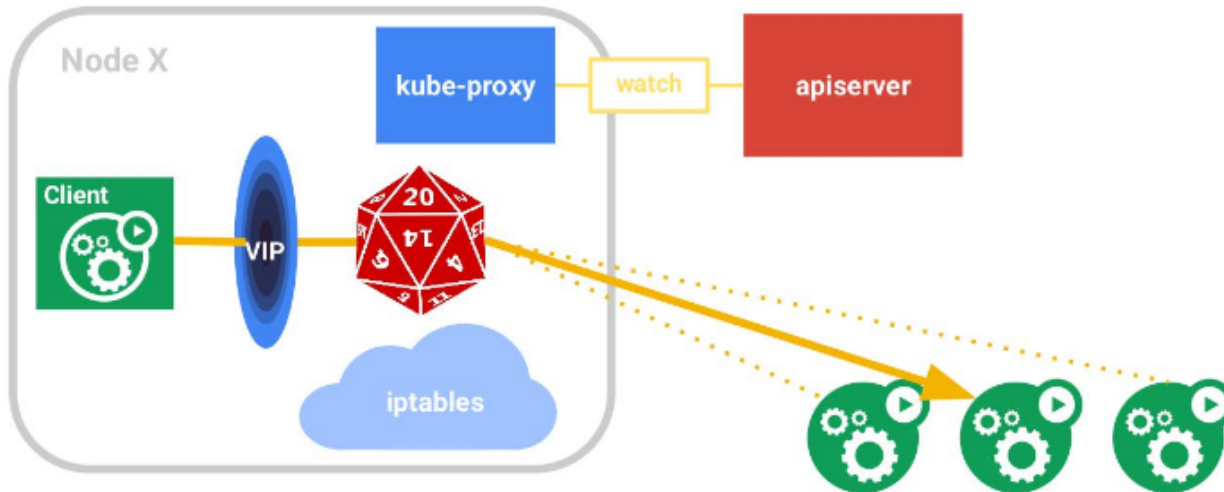
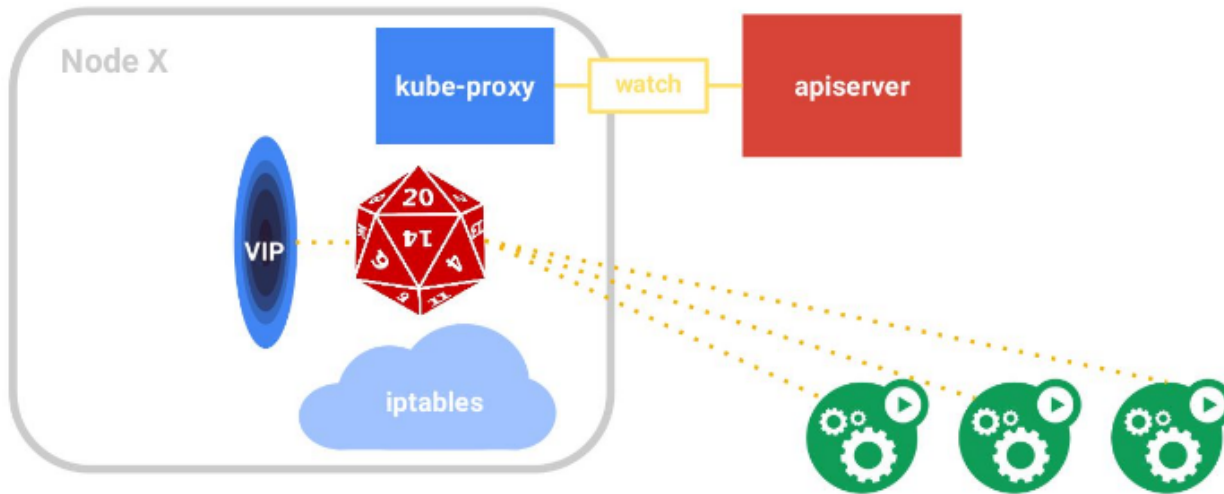
`kubectl run + kubectl expose`













Refs

Refs

1. Tim Hockin, Kubernetes: One Year Later
2. Ray Tsang @saturnism, Kubernetes with Java-based Microservices
3. Stefan Schimanski, Kubernetes Architecture & Introduction
4. @luxas - [kubernetes-on-arm](#)
5. Daniel Smith, What's new in Kubernetes
6. [Kubernetes - What is Kubernetes?](#)
7. [Kubernetes - User Guide](#)
8. [Learn the Kubernetes Key Concepts](#)
9. [Kubernetes Intro and Update @thockin](#)
10. [kubernetes/architecture.md at release-1.2 - kubernetes/kubernetes](#)
11. [Kubernetes - User Guide](#)
12. [An Introduction to Kubernetes](#)



END

Eueung Mulyana

<http://eueung.github.io/docker-stuff/kubernetes>

CodeLabs | Attribution-ShareAlike CC BY-SA